

Approximate Solution

Charles Kimbrough



Approach

1. Start With a Random Assignment
 - a. Each variable is randomly set to True/False
 - b. All clauses are evaluated once
 - c. A list for each of the satisfied and unsatisfied clauses is created
2. Repeatedly Fix an Unsatisfied Clause
 - a. Every iteration chooses one random unsatisfied clause and tries to fix it using one of two strategies. Randomly selected which one is chosen.

Strategy A: Random Walk

Has a probability of 0.4 to occur

Pick a literal from the unsatisfied clause and flip its variable

This is to escape any local minima by introducing random noise.

Strategy B: Greedy Improvement

Has a probability of 0.6 to occur

Evaluate the gain from flipping each variable in the clause

$$\text{gain}(x) = \Delta(\text{number of satisfied clauses})$$

Pick the variable with the best positive gain, tie-broken randomly

This is to climb uphill by making the single flip that increases satisfaction the most.

Approach Continued

After a certain number of max flips, the algorithm:

- saves its best score so far
- completely restarts with a new random assignment

This is to avoid getting stuck in a hole, and try a new selection

This runs for a certain time before quitting, to satisfy the anytime.

If the best score equals number of clauses, the algorithm quits early.

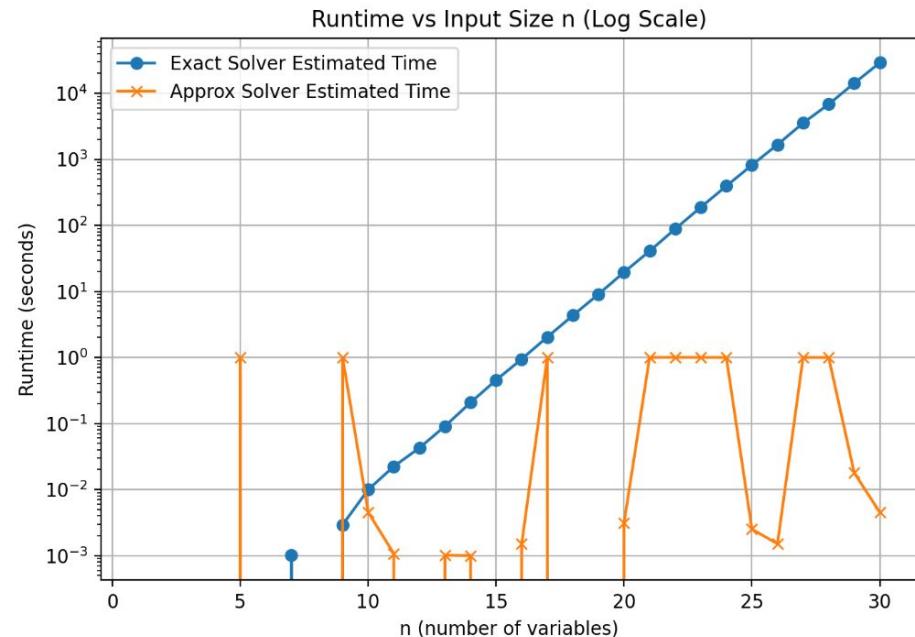
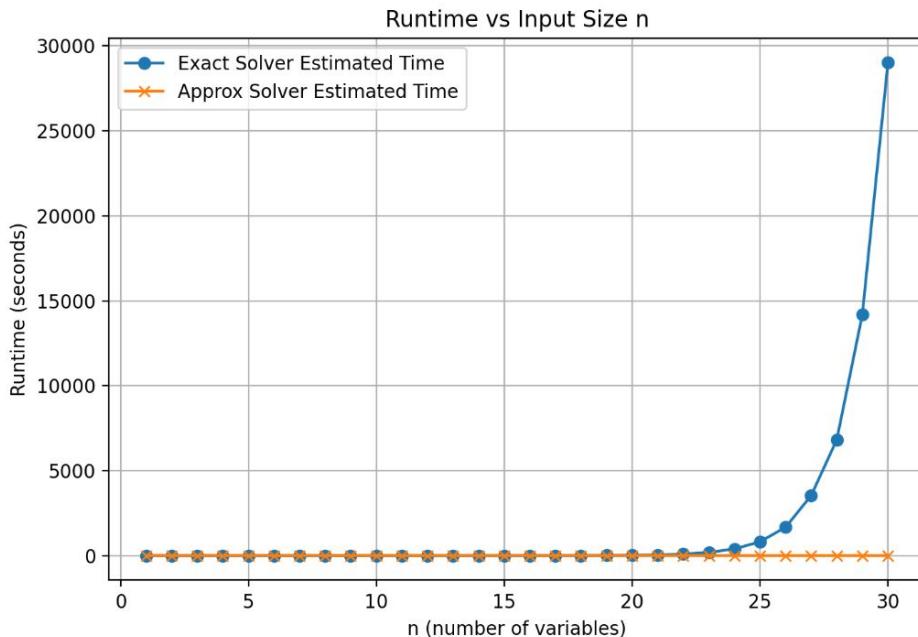
Runtime

n = number of variables

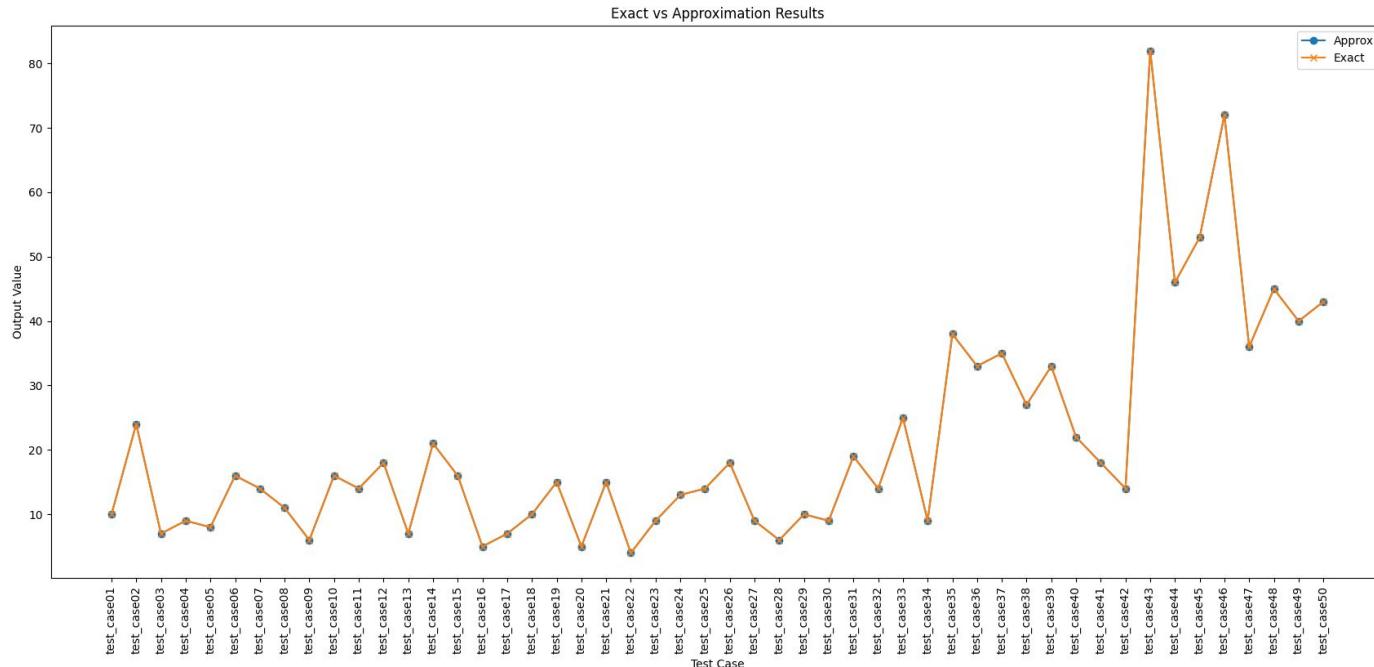
m = number of clauses

The cost per flip is $O(m/n)$. Since this runs for a constant number of flips, the runtime for a single restart is $O(m/n)$. Since most of time $m \gg n$, it is closer to $O(m)$.

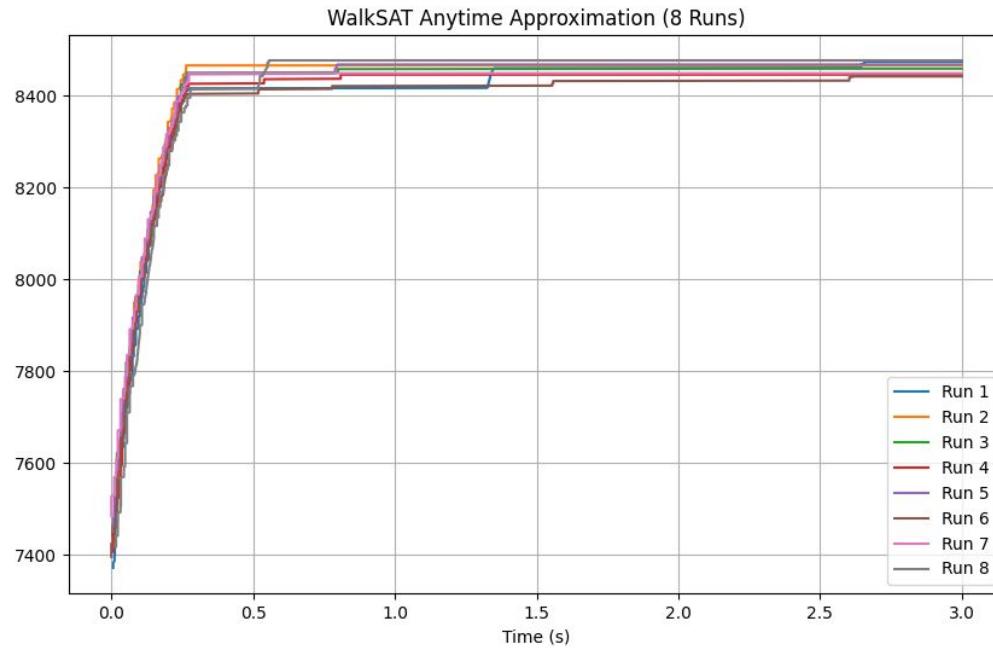
Approximation vs Exact Solution



Approximation vs Exact Solution



Solution Over Runtime



Parallelism

- Adding parallelism, decreased the initial speed of finding a solution.
- There was no consistent improvement in results for different parallelism.

